1. Advisor Feedback – Areas to Improve from Sprint 1

Area Improve 1 – LLM Integration Strategy

• Advisor comment:

Advisors recommended spending more effort defining the LLM integration approach and suggested using OpenAI within a small (<$30) budget to prototype AI functionality early.

• Why this matters:

The LLM is a core differentiator of the platform, enabling article simplification, explanations, and intelligent personalization. Without early validation, there is risk that AI latency, cost, or quality may not meet usability and performance expectations. Establishing a concrete, budget-aware LLM strategy reduces technical and financial uncertainty.

• Sprint 2 change(s) made:

- Began integrating an AI assistant for article explanations and simplification.
- Adopted a low-cost OpenAI usage approach for controlled experimentation.
- Implemented Hugging Face readability/complexity scoring to support adaptive content tailoring.
- Defined groundwork for future LLM-powered summarization and personalization features.

Area Improve 2 – Recommendation Logic Based on User Behavior

• Advisor comment:

Advisors questioned how article recommendations would be generated using users' reading history and interactions.

• Why this matters:

Effective recommendations are essential for personalization and user engagement. Without a clear strategy, recommendations risk being irrelevant, repetitive, or non-adaptive, which would reduce platform value and user retention.

• Sprint 2 change(s) made:

- Implemented topic filtering and user-driven content exploration.
- Introduced embeddings-based similarity groundwork to enable future recommendation ranking.
- Added feedback loop foundation for improving recommendations based on user interaction.

- Began designing logic to avoid showing duplicate or previously read articles.

Area Improve 3 – News Data Ingestion & Storage Strategy

• **Advisor comment:**

Advisors asked how news data is retrieved from APIs, when content should be saved, and how updates and duplicates are handled.

• **Why this matters:**

A clear ingestion strategy ensures data consistency, prevents duplicate content, and guarantees users receive fresh and relevant news. Without caching and update logic, the system may repeatedly show identical articles or miss new content.

• **Sprint 2 change(s) made:**

- Implemented article caching to avoid repeated API calls.
- Added logic to fetch updated articles per topic instead of repeating the same content.
- Improved deduplication mechanism to prevent duplicate articles appearing in feeds.
- Refined categorization and timestamping to track article freshness and ordering.

Area Improve 4 – System Administration & Data Management

• **Advisor comment:**

Advisors highlighted the need for clearer user and system management capabilities as the platform scales.

• **Why this matters:**

Administrative control is necessary for maintaining data integrity, managing user accounts, and ensuring safe operation as the system grows.

• **Sprint 2 change(s) made:**

- Implemented admin capabilities to delete and manage users.
- Improved backend support for categorization and timestamp management.
- Strengthened data lifecycle handling for stored and cached content.

Area Improve 5 – User Personalization & Content Adaptation

• **Advisor comment:**

Advisors encouraged clearer mechanisms for adapting content complexity and helping users progressively improve their understanding.

• **Why this matters:**

Personalization and adaptive readability are central to the platform's educational and user-experience value. Without structured scoring and progression, AI simplification may feel inconsistent or unhelpful.

• **Sprint 2 change(s) made:**

- Integrated Hugging Face readability scoring to evaluate article complexity.
- Built the foundation for AI-driven explanations and adaptive content simplification.
- Began implementing a user improvement roadmap, allowing the system to guide users toward progressively deeper understanding.

2. Peer Feedback on Sprint 1 Demo

- It would be helpful to more clearly show how user preferences (topics, complexity level) actually influence which articles are shown, especially when articles fit multiple categories.

• Peer Feedback Item 1 (paraphrased):

Show more clearly how user preferences and reading behavior influence which articles are recommended, especially when articles belong to multiple categories.

• **Consideration status:** Implemented (Foundation level)

• **Explanation:**

This feedback aligned directly with our platform's core goal of personalization. While full transparency indicators are not yet complete, Sprint 2 focused on strengthening the underlying recommendation logic and ensuring content is filtered and surfaced based on user-selected topics and history.

• **If implemented: related Sprint 2 change(s) or backlog item(s):**

- Topic filtering and improved categorization
- Embeddings groundwork for similarity-based recommendation
- Avoid showing duplicate or previously read articles
- Feedback loop foundation for improving recommendations

- As you build out more features, consider adding some personalization indicators (like "recommended for you" or match scores) so users can see why certain content is being surfaced.

• Peer Feedback Item 2 (paraphrased):

Add visible indicators such as "Recommended for you," match score, or explanations so users understand *why* an article is shown.

• **Consideration status:** Partially considered

• **Explanation:**

We recognized this as important for transparency and trust. In Sprint 2, we focused on building the recommendation logic and AI explanation infrastructure first. Visible explanation indicators are planned as a UI enhancement once recommendation confidence and scoring are fully integrated.

• **If implemented: related Sprint 2 change(s) or backlog item(s):**

- AI-driven explanations assistant groundwork
- Embeddings and recommendation scoring preparation

I thought the UX was very appealing, and the layout / functionality was outstanding. If I could, I would echo some previous comments regarding the "why" behind the recommended articles, as well as some more information about those articles (date, author, publisher) to help give the user a more informed picture of the link they are about to click. Solid work overall!

• Peer Feedback Item 3A (paraphrased):

Provide clearer transparency on *why* certain articles are recommended to the user.

• **Consideration status:** Partially considered

• **Explanation:**

This feedback aligns with improving trust and clarity in personalization. During Sprint 2, the team focused on strengthening the underlying recommendation infrastructure (topic filtering, embeddings groundwork, and feedback loop foundation). However, visible indicators explaining *why* an article is recommended (e.g., "Recommended for you," match score, or explanation) have not yet been implemented and remain planned for future refinement.

• **If implemented: related Sprint 2 change(s) or backlog item(s):**

- Topic filtering and improved categorization
- Embeddings groundwork for similarity-based recommendations
- Feedback loop foundation for improving recommendation quality

• Peer Feedback Item 3B (paraphrased):

Add more contextual information in article previews (e.g., publication date, author, publisher) to help users make informed decisions before clicking.

• **Consideration status:** Deferred / Not Implemented

• **Explanation:**

While the system processes article metadata internally, displaying richer contextual information in the article preview UI was not fully implemented in Sprint 2. This enhancement was deprioritized in favor of completing core platform functionality and AI-related improvements. The team recognizes its importance for usability and credibility and plans to address it in a future iteration if time permits.

• **If implemented: related Sprint 2 change(s) or backlog item(s):**

N/A — Not implemented in Sprint 2.

> One area worth exploring further is how your system addresses potential bias in scraped articles. Since coverage can differ widely across outlets, it could be valuable to let users create dedicated folders or collections for specific issues they're interested in. Within these, users could add articles from various sources they personally consider credible, allowing them to compare how different outlets frame the same topic. This kind of feature would encourage a more balanced, customizable news experience and help users develop a more holistic understanding of complex issues.

• Peer Feedback Item 4 (paraphrased):

Address potential bias and allow comparison across multiple sources for balanced perspectives.

• **Consideration status:** Deferred

• **Explanation:**

This is a valuable long-term enhancement but outside the core scope for the current sprint. The team prioritized core AI functionality and recommendation infrastructure first. Bias indicators or multi-source comparison may be explored in future iterations if time permits.

> Really like the idea behind this project. I think that a piece of feedback would be that if the user has differing expertise levels for different categories, the app should be able to reflect that. Right now it seems as though you can only select one expertise level for all for your interests.

• Peer Feedback Item 5 (paraphrased):

Support more flexible expertise or complexity levels (e.g., different levels per topic instead of global).

• **Consideration status:** Partially considered

• **Explanation:**

Sprint 2 introduced readability and complexity scoring using Hugging Face models to evaluate article difficulty, which lays the foundation for more granular adaptation. However, per-topic expertise adaptation has not yet been fully implemented.

• **If implemented: related Sprint 2 change(s) or backlog item(s):**

- Readability and complexity scoring integration
- AI-driven explanations groundwork